

CS 240 – Data Structures and Data Management

Module 5: Other Dictionary Implementations

Armin Jamshidpey

Based on lecture notes by many previous cs240 instructors

David R. Cheriton School of Computer Science, University of Waterloo

Winter 2023

version 2024-05-15 09:06

Outline

5 Dictionaries with Lists revisited

- Dictionary ADT: Implementations thus far
- Skip Lists
- Re-ordering Items

Outline

5 Dictionaries with Lists revisited

- Dictionary ADT: Implementations thus far
 - Skip Lists
 - Re-ordering Items

Dictionary ADT: Implementations thus far

A *dictionary* is a collection of key-value pairs (KVPs), supporting operations *search*, *insert*, and *delete*.

Realizations we have seen so far:

- Unordered array or linked list: $\Theta(1)$ insert, $\Theta(n)$ search and delete
- Ordered array: $\Theta(\log n)$ search, $\Theta(n)$ insert and delete
- Binary search trees: $\Theta(\text{height})$ search, insert and delete
- Balanced BST (AVL trees):
 $\Theta(\log n)$ search, insert, and delete

Dictionary ADT: Implementations thus far

A *dictionary* is a collection of key-value pairs (KVPs), supporting operations *search*, *insert*, and *delete*.

Realizations we have seen so far:

- **Unordered array or linked list:** $\Theta(1)$ insert, $\Theta(n)$ search and delete
- **Ordered array:** $\Theta(\log n)$ search, $\Theta(n)$ insert and delete
- **Binary search trees:** $\Theta(\text{height})$ search, insert and delete
- **Balanced BST** (AVL trees):
 $\Theta(\log n)$ search, insert, and delete

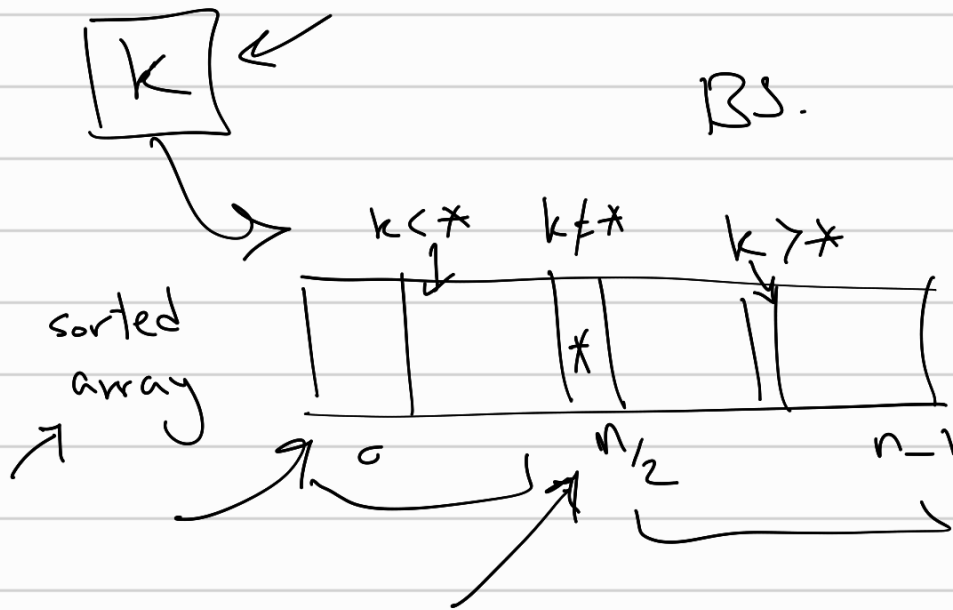
Improvements/Simplifications?

- **Can show:** If the KVPs were inserted in random order, then the expected height of the binary search tree would be $O(\log n)$.
- How can we use randomization within the data structure to mirror what would happen on random input?

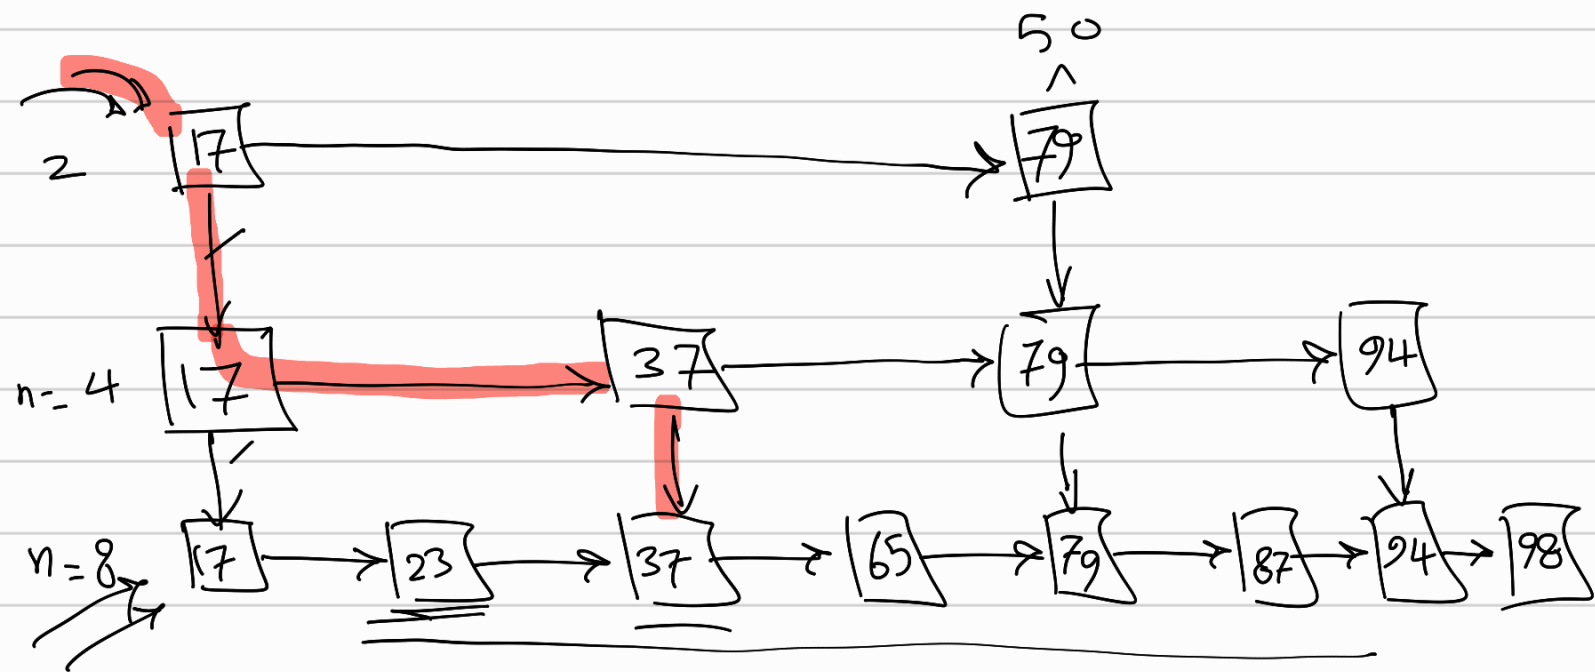
Outline

5 Dictionaries with Lists revisited

- Dictionary ADT: Implementations thus far
- Skip Lists
- Re-ordering Items



$$O(\log n)$$



Search(50) X

$$2 + 4 + 8 \leq 16$$

$$h = 2 \leq \log 8 = 3$$

$$h \in O(\log n)$$


$$2^{(n-1)} \leftarrow 2^{(2^t-1)}$$

$$n = 2^t$$

$$2^t + 2^{t-1} + 2^{t-2} + 2^{t-3} + \dots + 2 = 2(2^{t-1} + \dots + 1)$$


Skip Lists

- A hierarchy S of ordered linked lists (*levels*) S_0, S_1, \dots, S_h :
 - Each list S_i contains the special keys $-\infty$ and $+\infty$ (sentinels)
 - List S_0 contains the KVPs of S in non-decreasing order.
(The other lists store only keys, or links to nodes in S_0 .)
 - Each list is a subsequence of the previous one, i.e., $S_0 \supseteq S_1 \supseteq \dots \supseteq S_h$
 - List S_h contains only the sentinels; the left sentinel is the *root*



Skip Lists


- A hierarchy S of ordered linked lists (*levels*) S_0, S_1, \dots, S_h :
 - ▶ Each list S_i contains the special keys $-\infty$ and $+\infty$ (sentinels)
 - ▶ List S_0 contains the KVPs of S in non-decreasing order.
(The other lists store only keys, or links to nodes in S_0 .)
 - ▶ Each list is a subsequence of the previous one, i.e., $S_0 \supseteq S_1 \supseteq \dots \supseteq S_h$
 - ▶ List S_h contains only the sentinels; the left sentinel is the *root*



- Each KVP belongs to a tower of nodes
- There are (usually) more *nodes* than *keys*
- The skip list consists of a reference to the topmost left node.
- Each node p has references $p.after$ and $p.below$

Search in Skip Lists

For each level, find **predecessor** (node before where k would be).
This will also be useful for *insert/delete*.



getPredecessors (k)


1. $\rightarrow \underline{p} \leftarrow \text{root}$
2. $\rightarrow \underline{P} \leftarrow$ stack of nodes, initially containing \underline{p}
3. **while** $\underline{p}.\text{below} \neq \text{NIL}$ **do**
4. $\underline{p} \leftarrow \underline{p}.\text{below}$
5. **while** $\underline{p}.\text{after}.\text{key} \leq k$ **do** $\underline{p} \leftarrow \underline{p}.\text{after}$
6. $\underline{P}.\text{push}(\underline{p})$
7. **return** \underline{P}

skipList::search (k)

1. $\underline{P} \leftarrow \underline{\text{getPredecessors}}(k)$
2. $\underline{p_0} \leftarrow \underline{P}.\text{top}()$ // predecessor of \underline{k} in S_0
3. **if** $\underline{p_0}.\text{after}.\text{key} = k$ **return** $\underline{p_0}.\text{after}$
4. **else return** "not found, but would be after $\underline{p_0}$ "

Example: Search in Skip Lists

Example: *search*(87)




S

s_0	83
s_1	83
s_2	65
s_3	$-\infty$


Example: Search in Skip Lists

Example: *search*(87)



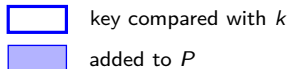

Example: Search in Skip Lists

Example: *search*(87)



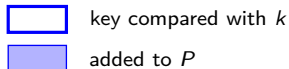

Example: Search in Skip Lists

Example: *search*(87)



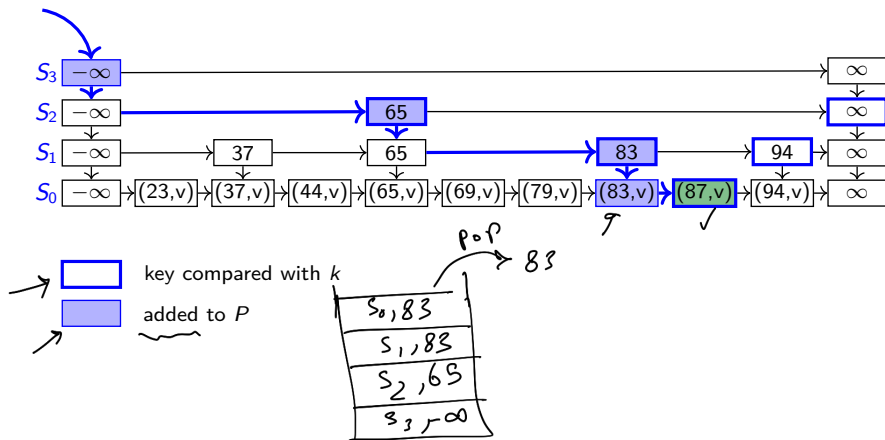
Example: Search in Skip Lists

Example: *search*(87)



Example: Search in Skip Lists

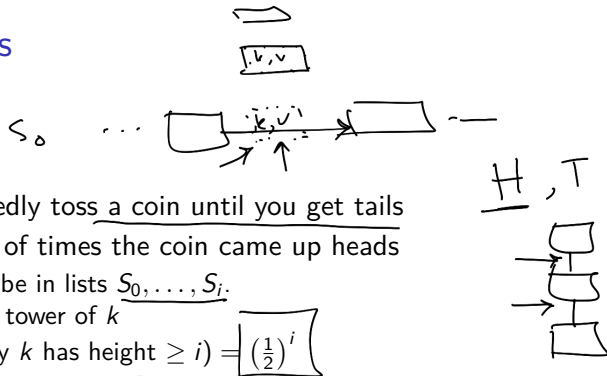
Example: *search*(87)



Insert in Skip Lists

skipList::insert(k, v)


- Randomly repeatedly toss a coin until you get tails
- Let i the number of times the coin came up heads
 - we want k to be in lists S_0, \dots, S_i .
 - $i \rightarrow$ **height** of tower of k
 - $P(\text{tower of key } k \text{ has height } \geq i) = \left(\frac{1}{2}\right)^i$
- Increase height h of skip list, if needed, to have $h > i$ levels.
- Use *getPredecessors*(k) to get stack P .
The top i items of P are the predecessors p_0, p_1, \dots, p_i of where k should be in each list S_0, S_1, \dots, S_i
- Insert (k, v) after p_0 in S_0 , and k after p_j in S_j for $1 \leq j \leq i$



Example: Insert in Skip Lists

Example: *skipList::insert*(52, *v*)

Coin tosses: H,T $\Rightarrow \underline{i = 1}$

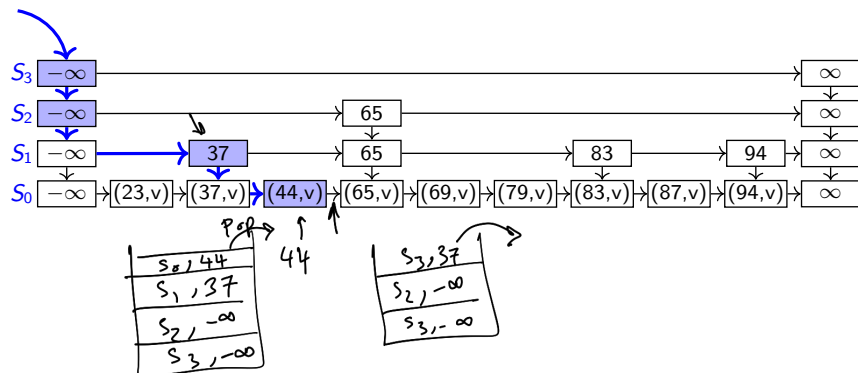


Example: Insert in Skip Lists

Example: *skipList::insert*(52, v)

Coin tosses: H,T $\Rightarrow i = 1$

getPredecessors(52)




Example: Insert in Skip Lists

Example: *skipList::insert*(52, v)

Coin tosses: H,T $\Rightarrow i = 1$


getPredecessors(52)



Example 2: Insert in Skip Lists

Example: *skipList::insert*(100, v)

Coin tosses: H,H,H,T $\Rightarrow i = 3$




Example 2: Insert in Skip Lists

Example: *skipList::insert*(100, v)

Coin tosses: H,H,H,T $\Rightarrow \underline{i = 3}$

Height increase




Example 2: Insert in Skip Lists

Example: *skipList::insert*(100, *v*)

Coin tosses: H,H,H,T $\Rightarrow i = 3$

Height increase

getPredecessors(100)




Example 2: Insert in Skip Lists

Example: *skipList::insert*(100, *v*)

Coin tosses: H,H,H,T $\Rightarrow i = 3$

Height increase

getPredecessors(100)



Insert in Skip Lists

skipList::insert(k, v)

```
1.  $\rightarrow P \leftarrow \text{getPredecessors}(k)$ 
2.  $\rightarrow$  for ( $i \leftarrow 0$ ;  $\text{random}(2) = 1$ ;  $i \leftarrow i+1$ ) {} // random tower height
3.   while  $i \geq P.\text{size}()$  // increase skip-list height?
4.    $\rightarrow$     $\left\{ \begin{array}{l} \text{root} \leftarrow \text{new sentinel-only list, linked in appropriately} \\ \text{add left sentinel of root at bottom of stack } P \end{array} \right.$ 
5.
6.    $p \leftarrow P.\text{pop}()$  // insert ( $k, v$ ) in  $S_0$ 
7.    $z_{\text{below}} \leftarrow \text{new node with } (k, v), \text{ inserted after } p$ 
8.   while  $i > 0$  // insert  $k$  in  $S_1, \dots, S_i$ 
9.      $p \leftarrow P.\text{pop}()$ 
10.     $z \leftarrow \text{new node with } k \text{ added after } p$ 
11.     $z.\text{below} \leftarrow z_{\text{below}}; z_{\text{below}} \leftarrow z$ 
12.     $i \leftarrow i - 1$ 
```

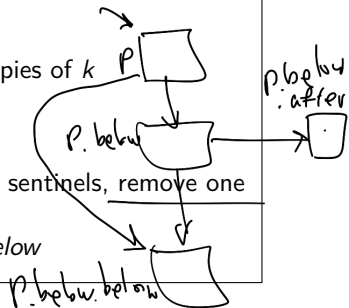

Delete in Skip Lists

It is easy to remove a key since we can find all predecessors.

Then eliminate layers if there are multiple ones with only sentinels.


skipList::delete(k)

1. $P \leftarrow \text{getPredecessors}(k)$
2. **while** P is non-empty
3. $p \leftarrow P.\text{pop}()$ // predecessor of k in some layer
4. **if** $p.\text{after}.key = k$
5. $p.\text{after} \leftarrow p.\text{after}.\text{after}$
6. **else break** // no more copies of k
7. $p \leftarrow$ left sentinel of the root-list
8. **while** $p.\text{below}.\text{after}$ is the ∞ -sentinel
 // the two top lists are both only sentinels, remove one
9. $p.\text{below} \leftarrow p.\text{below}.\text{below}$
10. $p.\text{after}.\text{below} \leftarrow p.\text{after}.\text{below}.\text{below}$



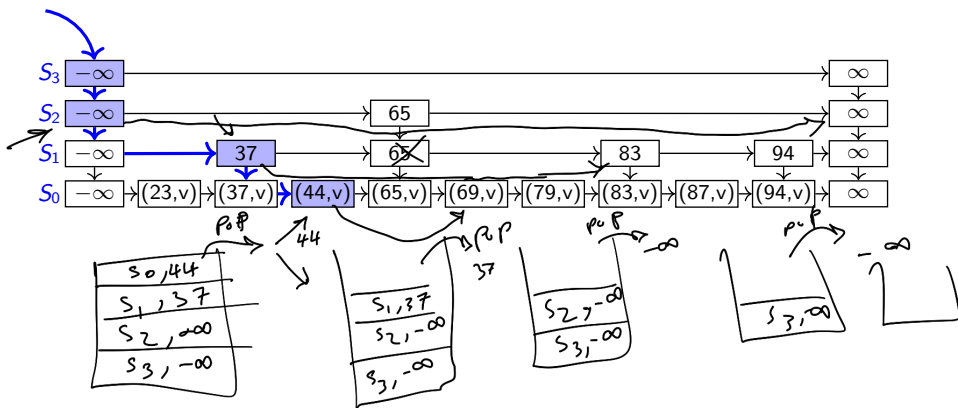
Example: Delete in Skip Lists

Example: *skipList::delete*(65)




Example: Delete in Skip Lists

Example: *skipList::delete*(65)
getPredecessors(65)



Example: Delete in Skip Lists

Example: *skipList::delete*(65)
getPredecessors(65)




Example: Delete in Skip Lists

Example: *skipList::delete*(65)

getPredecessors(65)

Height decrease




Analysis of Skip Lists



- Expected **space** usage: $O(n)$
- Expected **height**: $O(\log n)$
- Crucial for all operations:
 - ▶ How often do we **drop down** (execute $p \leftarrow p.\text{below}$)?
 - ▶ How often do we **step forward** (execute $p \leftarrow p.\text{after}$)?
- **skipList::search**: $O(\log n)$ expected time
 - ▶ # drop-downs = height
 - ▶ expected # forward-steps is ≤ 1 in each level
 - ▶ expected total # forward-steps is in $O(\log n)$
- **skipList::insert**: $O(\log n)$ expected time
- **skipList::delete**: $O(\log n)$ expected time

Summary of Skip Lists

- $O(n)$ expected space, all operations take $O(\log n)$ expected time.
- As described they are no faster than randomized binary search trees.
- Can show: A biased coin-flip to determine tower-height gives smaller expected run-times.
- Can save links (hence space) by implementing towers as array.



- Then skip lists are fast in practice and simple to implement.

Outline

5 Dictionaries with Lists revisited

- Dictionary ADT: Implementations thus far
- Skip Lists
- Re-ordering Items

Re-ordering Items

- Recall: Unordered list/array implementation of ADT Dictionary
search: $\Theta(n)$, insert: $\Theta(1)$, delete: $\Theta(1)$ (after a search)
- Lists/arrays are a very simple and popular implementation. Can we do something to make search more effective in practice?

Re-ordering Items

- Recall: Unordered list/array implementation of ADT Dictionary
search: $\Theta(n)$, *insert*: $\Theta(1)$, *delete*: $\Theta(1)$ (after a search)
- Lists/arrays are a very simple and popular implementation. Can we do something to make search more effective in practice?
- No: if items are accessed equally likely
- Yes: otherwise (we have a probability distribution of the items)
 - ▶ Intuition: Frequently accessed items should be in the front.
 - ▶ Two cases: Do we know the access distribution beforehand or not?
 - ▶ For short lists or extremely unbalanced distributions this may be faster than AVL trees or Skip Lists, and much easier to implement.

Optimal Static Ordering

Example:

	key	A	B	C	D	E
frequency of access		2	8	1	10	5
access-probability		$\frac{2}{26}$	$\frac{8}{26}$	$\frac{1}{26}$	$\frac{10}{26}$	$\frac{5}{26}$



- We count cost i for accessing the key in the i th position.

- Order A, B, C, D, E has expected access cost

$$\frac{2}{26} \cdot 1 + \frac{8}{26} \cdot 2 + \frac{1}{26} \cdot 3 + \frac{10}{26} \cdot 4 + \frac{5}{26} \cdot 5 = \frac{86}{26} \approx 3.31$$

- Order D, B, E, A, C has expected access cost

$$\frac{10}{26} \cdot 1 + \frac{8}{26} \cdot 2 + \frac{5}{26} \cdot 3 + \frac{2}{26} \cdot 4 + \frac{1}{26} \cdot 5 = \frac{66}{26} \approx 2.54$$



Optimal Static Ordering

Example:

key	A	B	C	D	E
frequency of access	2	8	1	10	5
access-probability	$\frac{2}{26}$	$\frac{8}{26}$	$\frac{1}{26}$	$\frac{10}{26}$	$\frac{5}{26}$


- We count cost i for accessing the key in the i th position.

- Order A, B, C, D, E has expected access cost

$$\frac{2}{26} \cdot 1 + \frac{8}{26} \cdot 2 + \frac{1}{26} \cdot 3 + \frac{10}{26} \cdot 4 + \frac{5}{26} \cdot 5 = \frac{86}{26} \approx 3.31$$

- Order D, B, E, A, C has expected access cost


$$\frac{10}{26} \cdot 1 + \frac{8}{26} \cdot 2 + \frac{5}{26} \cdot 3 + \frac{2}{26} \cdot 4 + \frac{1}{26} \cdot 5 = \frac{66}{26} \approx 2.54$$



- **Claim:** Over all possible static orderings, the one that sorts items by non-increasing access-probability minimizes the expected access cost.
- **Proof Idea:** For any other ordering, exchanging two items that are out-of-order according to their access probabilities makes the total cost decrease.

Dynamic Ordering: MTF


- What if we do *not know the access probabilities* ahead of time?
- Rule of thumb (**temporal locality**): A recently accessed item is likely to be used soon again.
- In list: Always insert at the front
- **Move-To-Front heuristic** (MTF): Upon a successful search, move the accessed item to the front of the list



- We can also do MTF on an array, but should then insert and search from the *back* so that we have room to grow.


Dynamic Ordering: Transpose

Transpose heuristic: Upon a successful search, swap the accessed item with the item immediately preceding it



Dynamic Ordering: Transpose

Transpose heuristic: Upon a successful search, swap the accessed item with the item immediately preceding it



Performance of dynamic ordering:

- Transpose does not adapt quickly to changing access patterns.
- MTF works well in practice.
- **Can show:** MTF is “2-competitive”:
No more than twice as bad as the optimal static ordering.